

```
In [0]: #This is a supplementary material to the lecture "Linear Regression" to quickly revise, whenever needed
```

```
In [0]: #Linear Regression tries to find a linear relation between the input features (say 'x') and the continuous spectrum output (say 'y')
#In other words, it tries to find the best possible line to fit the training data points
#after that, it uses the same linear function to predict the output for the unseen data
```

```
In [0]: #Linear regression will learn some coefficients and an intercept for the linear line and then, uses the same parameters to predict on new input values
```

```
In [0]: #Let's take an example of California housing dataset available in sklearn
```

```
In [0]: #import packages
import numpy as np
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [0]: #Load the data
housing = datasets.fetch_california_housing()
housing
```

Downloading Cal. housing from <https://ndownloader.figshare.com/files/5976036> to /root/scikit_learn_data

```
Out[0]: {'DESCR': '.. _california_housing_dataset:\n\nCalifornia Housing dataset\n-----\n\n**Data Set Characteristics:**\n\n :Number of Instances: 20640\n\n :Number of Attributes: 8\nnumeric, predictive attributes and the target\n\n :Attribute Information:\n      - MedInc\nmedian income in block\n      - HouseAge      median house age in block\n      - AveRooms      average number of rooms\n      - AveBedrms     average number of bedrooms\n      - Population    house block population\n      - AveOccup      average house occupancy\n      - Latitude      house block latitude\n      - Longitude     house block longitude\n\n :Missing Attribute Values: None\n\nThis dataset was obtained from the StatLib repository.\nhttp://lib.stat.cmu.edu/datasets/\n\nThe target variable is the median house value for California districts.\n\nThis dataset was derived from the 1990 U.S. census, using one row per census block group. A block group is the smallest geographical unit for which the U.S. Census Bureau publishes sample data (a block group typically has a population of 600 to 3,000 people).\n\nIt can be downloaded/loaded using the function: sklearn.datasets.fetch_california_housing\n\n.. topic:: References\n\n - Pace, R. Kelley and Ronald Barry, Sparse Spatial Autoregressions,\n  Statistics and Probability Letters, 33 (1997) 291-297\n',
  'data': array([[ 8.3252, 41., 6.98412698, ..., 2.55555556,
                  37.88, -122.23, ],
                 [ 8.3014, 21., 6.23813708, ..., 2.10984183,
                  37.86, -122.22, ],
                 [ 7.2574, 52., 8.28813559, ..., 2.80225989,
                  37.85, -122.24, ],
                 ...,
                 [ 1.7, 17., 5.20554273, ..., 2.3256351,
                  39.43, -121.22, ],
                 [ 1.8672, 18., 5.32951289, ..., 2.12320917,
                  39.43, -121.32, ],
                 [ 2.3886, 16., 5.25471698, ..., 2.61698113,
                  39.37, -121.24, ]]),
  'feature_names': ['MedInc',
                    'HouseAge',
                    'AveRooms',
                    'AveBedrms',
                    'Population',
                    'AveOccup',
                    'Latitude',
                    'Longitude'],
  'target': array([4.526, 3.585, 3.521, ..., 0.923, 0.847, 0.894])}
```

```
In [0]: x, y = housing['data'], housing['target']
        print(x.shape)
        print(y.shape)

(20640, 8)
(20640,)
```

```
In [0]: #so we have 20640 data points and each data point has 8 features and one output value, which is average house value in units of 100,000.
        #split it in train and test parts
        x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2, random_state = 0) #20% of the total data will be test data

        #random_state will ensure the same data goes to train and test each time you run the program
```

```
In [0]: lr = LinearRegression()
        lr.fit(x_train, y_train)
```

```
Out[0]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
In [0]: #Let's see coefficients and intercept of the linear function, it has fit the training data on
        print('coefficients:', lr.coef_)
        print('intercept: ', lr.intercept_)

coefficients: [ 4.33333407e-01  9.29324337e-03 -9.86433739e-02  5.93215487e-01
 -7.56192502e-06 -4.74516383e-03 -4.21449336e-01 -4.34166041e-01]
intercept:  -36.85856910680116
```

```
In [0]: #Let's predict for the test data i.e. unseen data
        y_pred = lr.predict(x_test)
```

```
In [0]: #Let's see score of our model
        score_test = lr.score(x_test, y_test)
        score_train = lr.score(x_train, y_train)
        print('Training score: ', score_train)
        print('Testing score: ', score_test)

Training score:  0.6088968118672871
Testing score:  0.5943232652466175
```

```
In [0]: #Thanks, happy Coding!
```

```
In [ ]: #To download .ipynb notebook, right click the following url and choose 'save link as'
        https://ninjasfiles.s3.amazonaws.com/0000000000003732.ipynb
```